Chris Ying

chrisying@cmu.edu

> EDUCATION

CARNEGIE MELLON UNIVERSITY (MLD)
MS in Machine Learning
2016 - 2017 | Pittsburgh, PA

CARNEGIE MELLON UNIVERSITY (SCS)
BS in Computer Science
Minor in Machine Learning
2013 - 2016 | Pittsburgh, PA
Cumulative GPA: 3.98/4.0

> CONTACT

Phone:// (408) 513-5636
Mail:// SMC 2302
Carnegie Mellon University
Pittsburgh, PA 15289
Github:// chrisying

> CONCENTRATIONS

Machine Learning
Natural Language Processing
Computer Vision
Parallel Architecture
Distributed Systems

> SKILLS

PROGRAMMING
Very experienced:
Python • C • Java • Go •
MATLAB/Octave • SML • LaTeX
Proficient:
C++ • Javascript
Familiar:
Assembly • MySQL
HTML • CSS • HQ9+

> AWARDS

2016 SCS High Honors
2016 Six-time Dean's List
2013 HackCMU Sponsor Winner
2013 Valedictorian of CHS

> POSITIONS

2016 President of ACM@CMU
2016 Team lead of AWAP
2016 TA for CMU 15-451
2015 Organizer of HackCMU
2015 Vice-president of ACM@CMU
2014 Board member of ACM@CMU

> EXPERIENCE

DROPBOX | Software Engineering Intern
May 2016 - Aug 2016 | San Francisco, CA

- Worked on rebuilding the internal logging pipeline using Apache Kafka.
- Shipped a tool for viewing logs in real-time that reduces latency from logging to viewing from hours to seconds.
- Coded primarily in Python and Javascript.

GOOGLE | Software Engineering Intern
May 2015 – Aug 2015 | Mountain View, CA

- Researched and prototyped various database options for a business reporting system including internal MySQL, NoSQL-like databases, and Cloud services.
- Built MapReduce jobs to process hundreds of gigabytes of customer dimensions and metrics.
- Designed a nearest neighbors algorithm for efficiently returning the peers of a customer in real-time.
- Coded primarily in Python, Go, and SQL-like languages.

GOOGLE | Software (Site Reliability) Engineering Intern
May 2014 – Aug 2014 | Seattle, CA

- Worked in a site reliability engineering team to write a library to streamline the process of monitoring internal services
- Automated the creation of graphs and alerts for various metrics.
- Coded primarily in Python and various configuration languages.

> PROJECTS

ML/NLP RESEARCH | Improving Event Coreference
Jan 2016 - May 2016
Designed and built an event coreference system which incorporates prior knowledge in order to better coreference event mentions in free text. The system parses events and entities from ECB+ data set, links it to existing knowledge bases (i.e. NELL, YAGO, DBpedia), and trains a logistic regression classifer for pairwise event coreference.

PARALLEL ARCHITECTURE | Parallel Binary Decision Diagrams
April 2016 - May 2016
Created a parallel binary decision diagram library for model checking which efficiently utilizes Intel Xeon Phi processors. The library was written from scratch using C++, Intel CilkPlus, and various lockfree data structures.

NATURAL LANGUAGE PROCESSING | Wikipedia Q&A System
Jan 2015 - May 2015
Utilized various NLP techniques to produce questions and answers from Wikipedia documents. The project was completed as a group of 4 and used Beautiful Soup, Stanford parser, dependency trees, coreference resolution, and logistic regression to parse and analyze raw HTML documents.

MACHINE LEARNING | CIFAR-10 Image Classifier
Nov 2014 - Dec 2014
Implemented a classifier using kernel multinomial logistic regression with principal component analysis to classify CIFAR-10, a 10-class image dataset. The system was optimized for performance to test the affects of PCA on runtime and accuracy.

AUDIENCE SENTIMENT WEB APP | Pulsr
Nov 2014 - Dec 2014
Developed a web application to track and distribute real time audience sentiment using dynamically generated graphs.